Practicum 2

Kieran Douglas

2. Examining Waugh's 1927 Asparagus Data

A.

We can see that the greatest difference between coeff estimates and those reported in the question are in the number of stalks and the variation in size (nostalks and disp). These end up being an absolute difference of 0.1767 and 0.0697 respecitvely.

```
# Set up environment and load data
library(tidyverse)
```

```
-- Attaching core tidyverse packages ------------------------ tidyverse 2.0.0 --
v dplyr     1.1.4     v readr     2.1.5
v forcats   1.0.0     v stringr   1.5.1
v ggplot2   3.5.1     v tibble    3.2.1
v lubridate 1.9.4     v tidyr     1.3.1
v purrr     1.0.2
-- Conflicts ------------------------------------------ tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()    masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to becor
```

```
library(readr)
library(fixest)

waugh <- read_table("/Users/kieran/Documents/MASTERS/METRICS/code/metrics/practicum_2_files/w
```

```
-- Column specification --------------------------------------------------------
cols(
```

```
  GREEN = col_double(),
  NOSTALKS = col_double(),
  DISPERSE = col_double(),
  PRICE = col_double()
)

Warning: 200 parsing failures.
row col  expected    actual
  1  -- 4 columns 5 columns '/Users/kieran/Documents/MASTERS/METRICS/code/metrics/practicum_1
  2  -- 4 columns 5 columns '/Users/kieran/Documents/MASTERS/METRICS/code/metrics/practicum_1
  3  -- 4 columns 5 columns '/Users/kieran/Documents/MASTERS/METRICS/code/metrics/practicum_1
  4  -- 4 columns 5 columns '/Users/kieran/Documents/MASTERS/METRICS/code/metrics/practicum_1
  5  -- 4 columns 5 columns '/Users/kieran/Documents/MASTERS/METRICS/code/metrics/practicum_1
... ... ......... ......... ..........................................................
See problems(...) for more details.
```

```
waugh_clean <- waugh |>
  rename(
    green = "GREEN",
    nostalks = "NOSTALKS",
    disp = "DISPERSE",
    price = "PRICE"
  ) |>
   mutate(
      green = as.numeric(green),
      nostalks = as.numeric(nostalks),
      disp = as.numeric(disp),
      price = as.numeric(price)
    )

# Run a MLR
model1 = lm(price ~ green+nostalks+disp, data = waugh_clean)
summary(model1)
```

```
Call:
lm(formula = price ~ green + nostalks + disp, data = waugh_clean)

Residuals:
    Min      1Q  Median      3Q     Max
-56.004  -9.485  -0.122   9.422  49.097
```

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) 40.761264   5.327837   7.651 8.82e-13 ***
green        0.137598   0.007099  19.382  < 2e-16 ***
nostalks    -1.357256   0.150822  -8.999  < 2e-16 ***
disp        -0.345283   0.129656  -2.663  0.00839 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 15.52 on 196 degrees of freedom
Multiple R-squared:  0.7268,    Adjusted R-squared:  0.7226
F-statistic: 173.8 on 3 and 196 DF,  p-value: < 2.2e-16
```

```r
# compare coefficient estimates
coeffs = coefficients(model1)
original_coeffs = c(green = 0.13826, nostalks = -1.53394, disp = -0.27554)
differences <- coeffs[names(original_coeffs)] - original_coeffs
print(differences)
```

```
        green      nostalks          disp
-0.0006617547  0.1766836397 -0.0697428425
```

```r
# We can see that the greatest difference between coeff estimates and those reported in the
```

B.

It looks like the main issue here is that Waugh has transformed green color on the asparagus variable to inches rather than what I have in the raw data, being hundredths of inches. To fix this I need to rescale the variable green by dividing it by 100. After doing this the difference in means is no longer a problem.

```r
means <- c(price_avg = mean(waugh_clean$price), green_avg = mean(waugh_clean$green), nostalks
print(means)
```

```
   price_avg    green_avg nostalks_avg     disp_avg
      90.095      588.750       19.555       14.875
```

```r
means_reported <- c(price_avg = 90.095, green_avg = 5.8875, nostalks_avg = 19.555, disp_avg =
differences_avg <- means[names(means_reported)] - means_reported
print(differences_avg)
```

```
    price_avg     green_avg  nostalks_avg      disp_avg
 0.000000e+00  5.828625e+02 -3.552714e-15  0.000000e+00
```

```
# It looks like the main issue here is that Waugh has transformed green color on the asparag
waugh_clean <- waugh_clean |>
  mutate(
    green = green/100
  )
# Recheck mean differences
means_recheck <- c(price_avg = mean(waugh_clean$price), green_avg = mean(waugh_clean$green),
print(means_recheck)
```

```
   price_avg    green_avg nostalks_avg     disp_avg
     90.0950       5.8875      19.5550      14.8750
```

C.

I notice relative variance in the relative size of the covariances reported. For example, the covariance between price and green according to Waugh is ~3430 while the one I found was 3448 (larger). On the other hand, Waugh found a covairance of ~-154 for green and disp while I found one of ~-180 (mine was smaller). In other cases though mine was larger like with the cov between green and price. I think that the pattern in differences may relate to the way my price variable is coded relative to Waugh's. This would explain the constant differences in price covariance versus the more consistent findings in ohter categories.

```
# I am now going to create a variance covariance matrix similar to Waugh's`
moments_subset <- waugh_clean[, c("price", "green", "nostalks", "disp")]
# to match the layout of the table provided I will modify green to be back to hundredths of a
moments_subset <- moments_subset |>
  mutate(
    green = green*100
  )
# Create the matric, format it correctly, and print
matrix <- cov(moments_subset)
print(matrix)
```

```
              price       green  nostalks       disp
price     868.73967  3448.18467 -93.38465  -87.43028
green    3448.18467 24439.38442 -17.09171 -180.05653
nostalks  -93.38465   -17.09171  60.73063   24.92399
disp      -87.43028  -180.05653  24.92399   83.48681
```

```
matrix_fmt <- formatC(matrix, format="f", digits=1)
matrix_fmt[lower.tri(matrix_fmt)] <- ""
print(noquote(matrix_fmt))
```

```
         price green   nostalks disp
price    868.7 3448.2  -93.4    -87.4
green          24439.4 -17.1    -180.1
nostalks               60.7     24.9
disp                            83.5
```

D.

It is plausible that the principle findings concerning the effects of covariate variation could be affected in terms of the precision and statistical inference of coefficient estimates if there was underestimation or misreporting of variance and covariance among regressors. Depsite OLS estimates remaining unbiased under GM, Waugh's under or over-estiamted variance and coveriance could lead to underestiated standard errors and possibly inacurate significance. Similarly, cov among the regressors included may inflate the standard errors and lead to reductions in the validity/reliability of estimates. The difference in scaled coefficients between Waugh's and my own analysis do not differ too much, but there are some clear differences. For example, my estimate for the marginal effect of a unit increase in the number of stalks per bunch is about \$0.5 less than Waugh's. I estimate a slightly higher increase in price given a unit increase in green, and a roughly \$0.2 larger decrease in price given a unit increase in size variation. Overall, our estimates are very similar. I find the association between both green and nostalks with price to be highly statistically significant at the $p<0.001$ level and the association between disp and price to be statistically signifiacnt at **$0.001<p<0.01$** level. All of these observations are ceterus parabus.

::: {.cell}

```
 # run model
coef(model1) * 2.782
```

::: {.cell-output .cell-output-stdout}

```
(Intercept)       green    nostalks        disp
113.3978351   0.3827983  -3.7758872  -0.9605769
```

:::

```
waugh_scaled <- original_coeffs*2.782
print(waugh_scaled)
```

::: {.cell-output .cell-output-stdout}

```
     green   nostalks       disp
 0.3846393 -4.2674211 -0.7665523
```

:::

```
# The difference in scaled coefficients between Waugh's and my own analysis do not differ too
```

:::

E.

To solve this I am going to experiment with a matrix.

```
# my matrix solutions
# Extract covariance matrix of regressors (columns and rows 2 to 4)
xx <- matrix[2:4, 2:4]
# Extract covariance vector of PRICE with regressors (row 1, columns 2 to 4)
xy <- matrix[1, 2:4]
# Compute beta_hat = solve(Sigma_XX) %*% Sigma_XY
beta_hat <- solve(xx) %*% xy
print(beta_hat)
```

```
              [,1]
green     0.1375982
nostalks -1.3572564
disp     -0.3452828
```

```
# Waugh's matrix solutions
vc_mat <- matrix(c(
  1063.64, 3430.89, -100.92, -82.35,
  3430.89, 24317.19, -17.01, -154.54,
  -100.92, -17.01, 61.33, 25.51,
  -82.35, -154.54, 25.51, 83.07
), nrow = 4, byrow = TRUE)
```

```r
# names for reference
row_col_names <- c("price","green","nostalks","disp")
dimnames(vc_mat) <- list(row_col_names, row_col_names)
# price-cov covariance vector
Sigma_yX <- vc_mat["price", c("green", "nostalks", "disp")]
# covariate-covariance matrix
Sigma_XX <- vc_mat[c("green", "nostalks", "disp"), c("green", "nostalks", "disp")]
# raw coefficients from matrix
beta_raw <- solve(Sigma_XX, Sigma_yX)
names(beta_raw) <- c("green", "nostalks", "disperse")
# scale
beta_final <- beta_raw * 2.782
beta_final
```

```
     green   nostalks   disperse
 0.3847304 -4.1520749 -0.7670883
```

```r
# Computing the least squares estimates from Waugh's VC matrix table, we can see that his est
```

3. Exploring Relationships among R^2, Coefficients of Determination, and Correlation Coefficients

A.

Based on the cor matrix, it is evident that the most highly correlated variables are price and green, price and nostalks, nostalks and disp, and price and disp (ordered from highest to lowest positive and negative correlation). The most orthogonal correlations are between green and nostalks and green and disp becasue each of their correlations is quite close to zero, thus implying low to no real linear relationship between the variables based on the data we have available.

```r
# create cor matrix
order <- c("price", "green", "nostalks", "disp")
cor_matrix <- cor(waugh_clean)
cor_matrix <- cor_matrix[order, order]

# round the matrix
cor_matrix_rounded <- round(cor_matrix, 5)
```

```r
# lower triangle with empty strings
cor_matrix_char <- format(cor_matrix_rounded, nsmall = 5)
cor_matrix_char[lower.tri(cor_matrix_char)] <- ""

print(cor_matrix_char, quote = FALSE)
```

```
         price    green    nostalks disp
price     1.00000  0.74834 -0.40656 -0.32464
green              1.00000 -0.01403 -0.12605
nostalks                    1.00000  0.35003
disp                                 1.00000
```

```r
# based on the cor matrix, it is evident that the most highly correlated variables are price
```

B.

Comparing the square roots calculated of the R^2 I can see that they are equal except for sign. This is because the correlation coefficient can be negative (between -1 and 1) but the coiefficient of determination will always be positive (between 0 and 1) since its the squared correlation coefficient. If I had run the reverse regressios the R^2 measures would be the same as those from the correct regressions because when dealing with the same two variables, the ammount of variation explained in one variable by the other will be constant since again, it is the squares correlation coefficient. Reversing their order has no impact on the ammount of variation explained that our R^2 captures.

```r
# price model
pricegreenmodel <- lm(data = waugh_clean, price ~ green)
pricenostalksmodel <- lm(data = waugh_clean, price ~ nostalks)
pricedispmodel <- lm(data = waugh_clean, price ~ disp)
#price on green
summary(pricegreenmodel)
```

```
Call:
lm(formula = price ~ green, data = waugh_clean)

Residuals:
   Min     1Q Median     3Q    Max
-64.85 -11.52   1.29  12.55  47.43
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)   7.0275     5.4130   1.298    0.196
green        14.1091     0.8888  15.875   <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 19.6 on 198 degrees of freedom
Multiple R-squared:   0.56, Adjusted R-squared:  0.5578
F-statistic:   252 on 1 and 198 DF,  p-value: < 2.2e-16
```

```
sqrt(0.56)
```

```
[1] 0.7483315
```

```
# price on nostalks
summary(pricenostalksmodel)
```

```
Call:
lm(formula = price ~ nostalks, data = waugh_clean)

Residuals:
    Min      1Q  Median      3Q     Max
-58.948 -17.317  -3.104   9.538  93.589

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 120.1645     5.1676  23.253  < 2e-16 ***
nostalks     -1.5377     0.2456  -6.262 2.32e-09 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 27 on 198 degrees of freedom
Multiple R-squared:  0.1653,    Adjusted R-squared:  0.1611
F-statistic: 39.21 on 1 and 198 DF,  p-value: 2.32e-09
```

```
sqrt(0.1653)
```

```
[1] 0.406571
```

```r
# price on disp
summary(pricedispmodel)
```

```
Call:
lm(formula = price ~ disp, data = waugh_clean)

Residuals:
    Min      1Q  Median      3Q     Max
-57.153 -19.295  -3.114  15.451  86.272

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 105.6726     3.7826   27.94  < 2e-16 ***
disp         -1.0472     0.2168   -4.83 2.73e-06 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 27.95 on 198 degrees of freedom
Multiple R-squared:  0.1054,    Adjusted R-squared:  0.1009
F-statistic: 23.33 on 1 and 198 DF,  p-value: 2.731e-06
```

```r
sqrt(0.1054)
```

```
[1] 0.3246537
```

C.

I expect that if I add the regressor nostalks to the regression, the R^2 will increase. This is because we have seperately confirmed that variation in nostalks explains some of the variation in price, so if we add it to the regression, the models explanatory power should increase. Given the correlation between green and nostalks in the table I would expect the change in R^2 when nostalks is added to the regression to be large. My intuition is such that since the two variables are not super correlated BUT are indidivually correlated with price, they should explain a significant ammount of the variation in price when both included as regressors. Running the regression confirms my intuition, as the R^2 is increased by nearly 15 points, representing a significant jump. Further, comparing the R^2 between the model that regressed price on green to the model that regressed price on green AND disp, we can see another increase in the R^2 equating to about 5 points. This represents an interesting jump, despite it being smaller in magnitude to the previous change in R^2. I think this is consistent with the sample correlation

between green and disp, which is bigger than that between green and nostalks but still not so big as to fail in adding additional explanation to the model. Finally, when comparing the price on nostalks model to the price on nostalks AND disp model, we also observe a jump in R^2 that is about 4 points. This is again smaller than previous jumps and can probably be explained by the higher correlation coefficient between nostalks and disp that comes out to 0.35003, an R that would suggest higher correlation between the two explanatory variables and thus a smaller increase to the R^2 when both included in a model explaining price. This is consistent with the previous findings regarding the effect of adding correlated explanatory variables to a model. I think this is why would prioritize using adjusted R^2 since that measure penalizes for adding less relavent variables.

```
green_nostalks_model <- lm(data = waugh_clean, price~green+nostalks)
summary(green_nostalks_model)
```

```
Call:
lm(formula = price ~ green + nostalks, data = waugh_clean)

Residuals:
    Min      1Q  Median      3Q     Max
-54.507  -9.997   0.746  10.008  50.997

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  36.9431     5.2100   7.091 2.33e-11 ***
green        14.0043     0.7148  19.593  < 2e-16 ***
nostalks     -1.4983     0.1434 -10.449  < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 15.76 on 197 degrees of freedom
Multiple R-squared:  0.7169,    Adjusted R-squared:  0.714
F-statistic: 249.5 on 2 and 197 DF,  p-value: < 2.2e-16
```

```
summary(pricegreenmodel)
```

```
Call:
lm(formula = price ~ green, data = waugh_clean)

Residuals:
```

```
   Min     1Q Median     3Q    Max
-64.85 -11.52   1.29  12.55  47.43

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)   7.0275     5.4130   1.298    0.196
green        14.1091     0.8888  15.875   <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 19.6 on 198 degrees of freedom
Multiple R-squared:   0.56, Adjusted R-squared:  0.5578
F-statistic:   252 on 1 and 198 DF,  p-value: < 2.2e-16
```

```
green_disp_model <- lm(data = waugh_clean, price~green+disp)
summary(green_disp_model)
```

```

Call:
lm(formula = price ~ green + disp, data = waugh_clean)

Residuals:
    Min      1Q  Median      3Q     Max
-64.610  -9.949   1.301  10.767  46.273

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  21.5318     5.7871   3.721 0.000259 ***
green        13.5529     0.8414  16.108  < 2e-16 ***
disp         -0.7549     0.1440  -5.244 4.03e-07 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 18.41 on 197 degrees of freedom
Multiple R-squared:  0.6139,    Adjusted R-squared:   0.61
F-statistic: 156.6 on 2 and 197 DF,  p-value: < 2.2e-16
```

```
summary(pricegreenmodel)
```

```

Call:
```

```
lm(formula = price ~ green, data = waugh_clean)

Residuals:
   Min     1Q Median     3Q    Max
-64.85 -11.52   1.29  12.55  47.43

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)   7.0275     5.4130   1.298    0.196
green        14.1091     0.8888  15.875   <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 19.6 on 198 degrees of freedom
Multiple R-squared:   0.56, Adjusted R-squared:  0.5578
F-statistic:   252 on 1 and 198 DF,  p-value: < 2.2e-16
```

```
nostalks_disp_model <- lm(data = waugh_clean, price~nostalks+disp)
summary(nostalks_disp_model)
```

```
Call:
lm(formula = price ~ nostalks + disp, data = waugh_clean)

Residuals:
    Min      1Q  Median      3Q     Max
-50.080 -16.153  -3.619  10.968  88.859

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 124.7557     5.2794  23.631  < 2e-16 ***
nostalks     -1.2626     0.2568  -4.917 1.85e-06 ***
disp         -0.6703     0.2190  -3.061  0.00252 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 26.44 on 197 degrees of freedom
Multiple R-squared:  0.2032,    Adjusted R-squared:  0.1951
F-statistic: 25.12 on 2 and 197 DF,  p-value: 1.923e-10
```

```r
summary(pricenostalksmodel)
```

```
Call:
lm(formula = price ~ nostalks, data = waugh_clean)

Residuals:
    Min      1Q  Median      3Q     Max
-58.948 -17.317  -3.104   9.538  93.589

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 120.1645     5.1676  23.253  < 2e-16 ***
nostalks     -1.5377     0.2456  -6.262 2.32e-09 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 27 on 198 degrees of freedom
Multiple R-squared:  0.1653,    Adjusted R-squared:  0.1611
F-statistic: 39.21 on 1 and 198 DF,  p-value: 2.32e-09
```

D.

Here I notice the same trend, where the R^2 for the full regression is 0.7268 while the summed R^2 from the single regressions is 0.8307. I think this is because in this case, the model accounts for overlaps in explanatory power between the regressors that are not even a question in the single models. I think this is becasue the multiple regression avoids double counting, providing the joint effect on price.

```r
fullmodel <- lm(data = waugh_clean, price~nostalks+green+disp)
summary(fullmodel)
```

```
Call:
lm(formula = price ~ nostalks + green + disp, data = waugh_clean)

Residuals:
    Min      1Q  Median      3Q     Max
-56.004  -9.485  -0.122   9.422  49.097
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  40.7613     5.3278   7.651 8.82e-13 ***
nostalks     -1.3573     0.1508  -8.999  < 2e-16 ***
green        13.7598     0.7099  19.382  < 2e-16 ***
disp         -0.3453     0.1297  -2.663  0.00839 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 15.52 on 196 degrees of freedom
Multiple R-squared:  0.7268,    Adjusted R-squared:  0.7226
F-statistic: 173.8 on 3 and 196 DF,  p-value: < 2.2e-16
```

```r
# R^2 of 0.7268
0.56+0.1653+0.1054
```

```
[1] 0.8307
```

```r
# sum of individual R^2 = 0.8307
```

E.

```r
# nostalks
numerator1 <- cov(waugh_clean$nostalks, waugh_clean$price) * (nrow(waugh_clean) - 1)
denominator1 <- var(waugh_clean$price) * (nrow(waugh_clean) - 1)
ratio1 <- numerator1 / denominator1
ratio1*-1.3573
```

```
[1] 0.1459021
```

```r
# green
numerator2 <- cov(waugh_clean$green, waugh_clean$price) * (nrow(waugh_clean) - 1)
denominator2 <- var(waugh_clean$price) * (nrow(waugh_clean) - 1)
ratio2 <- numerator2 / denominator2
ratio2*13.7598
```

```
[1] 0.5461513
```

```r
# disp
numerator3 <- cov(waugh_clean$disp, waugh_clean$price) * (nrow(waugh_clean) - 1)
denominator3 <- var(waugh_clean$price) * (nrow(waugh_clean) - 1)
ratio3 <- numerator3 / denominator3
ratio3*-0.3453
```

```
[1] 0.03475112
```